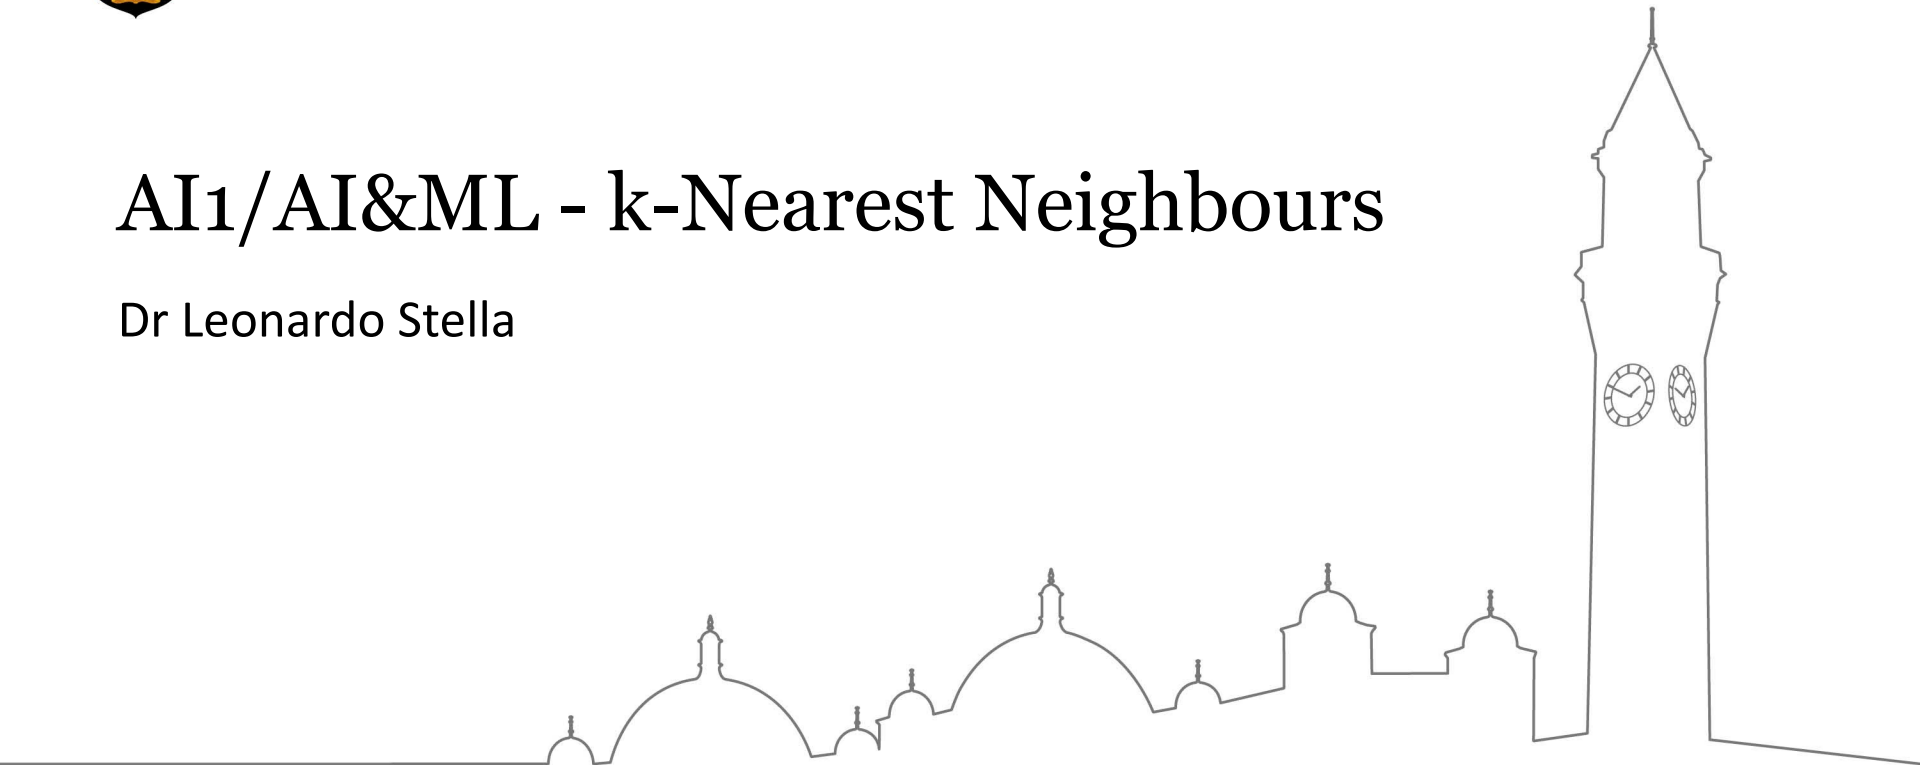




UNIVERSITY OF
BIRMINGHAM

AI1/AI&ML - k-Nearest Neighbours

Dr Leonardo Stella



Aims of the Session

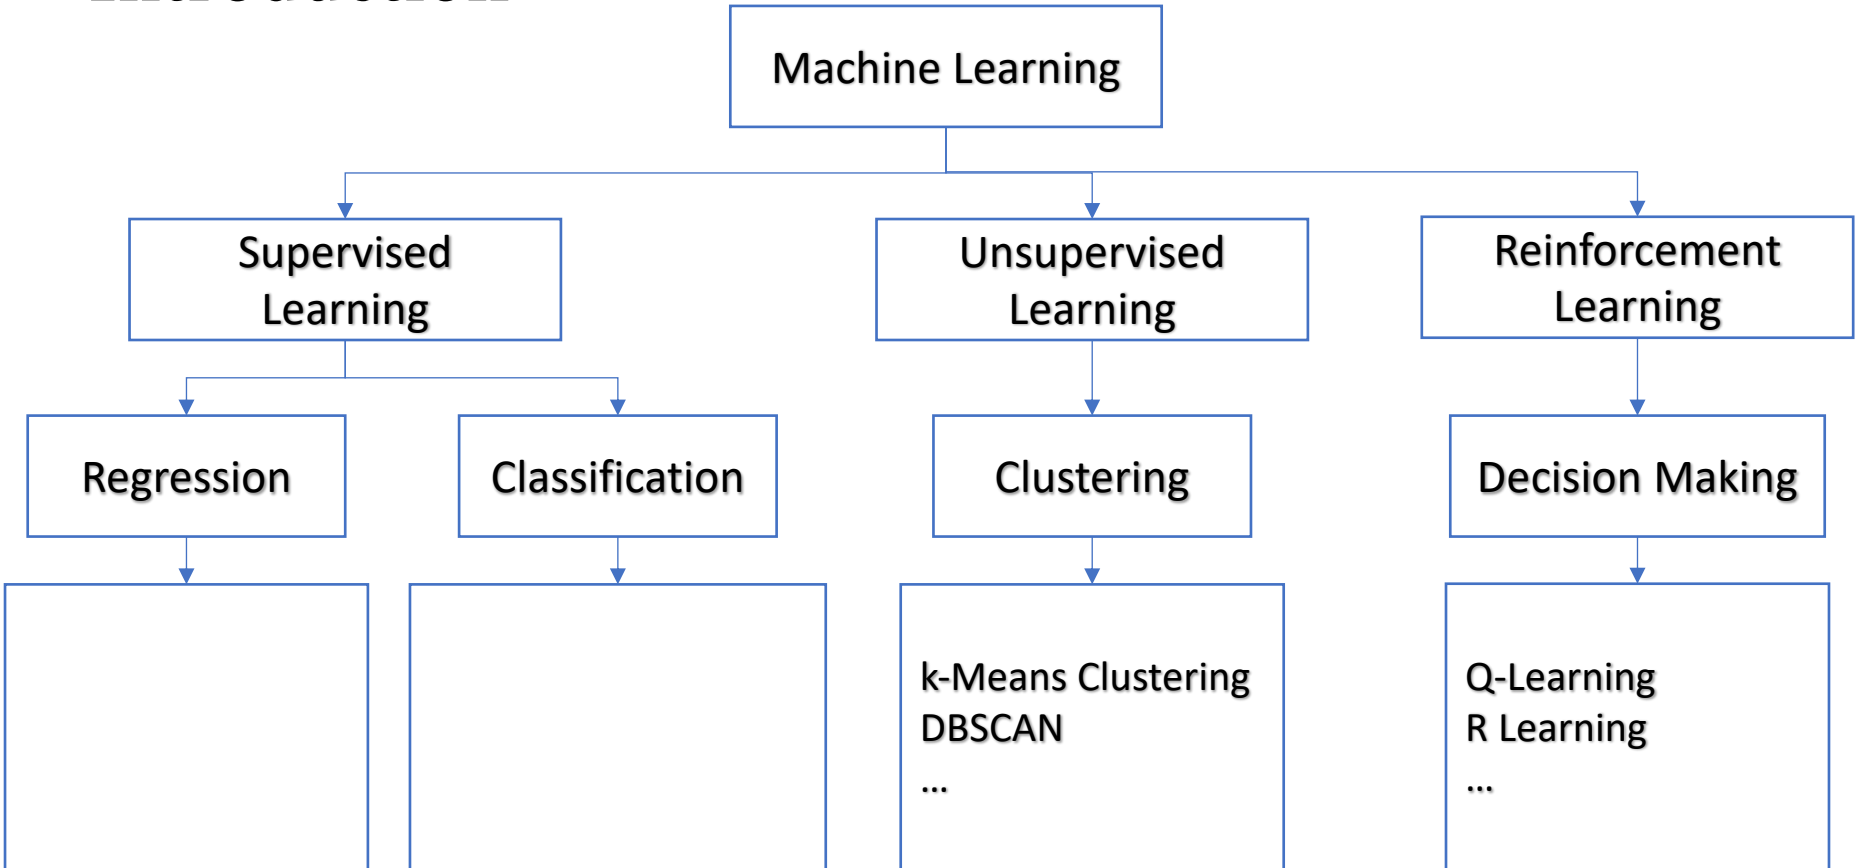
This session aims to help you:

- Explain the steps of k-Nearest Neighbours
- Apply k-Nearest Neighbours to problems involving numeric, ordinal and categorical input attributes
- Use evaluation procedures to estimate the performance of a predictor

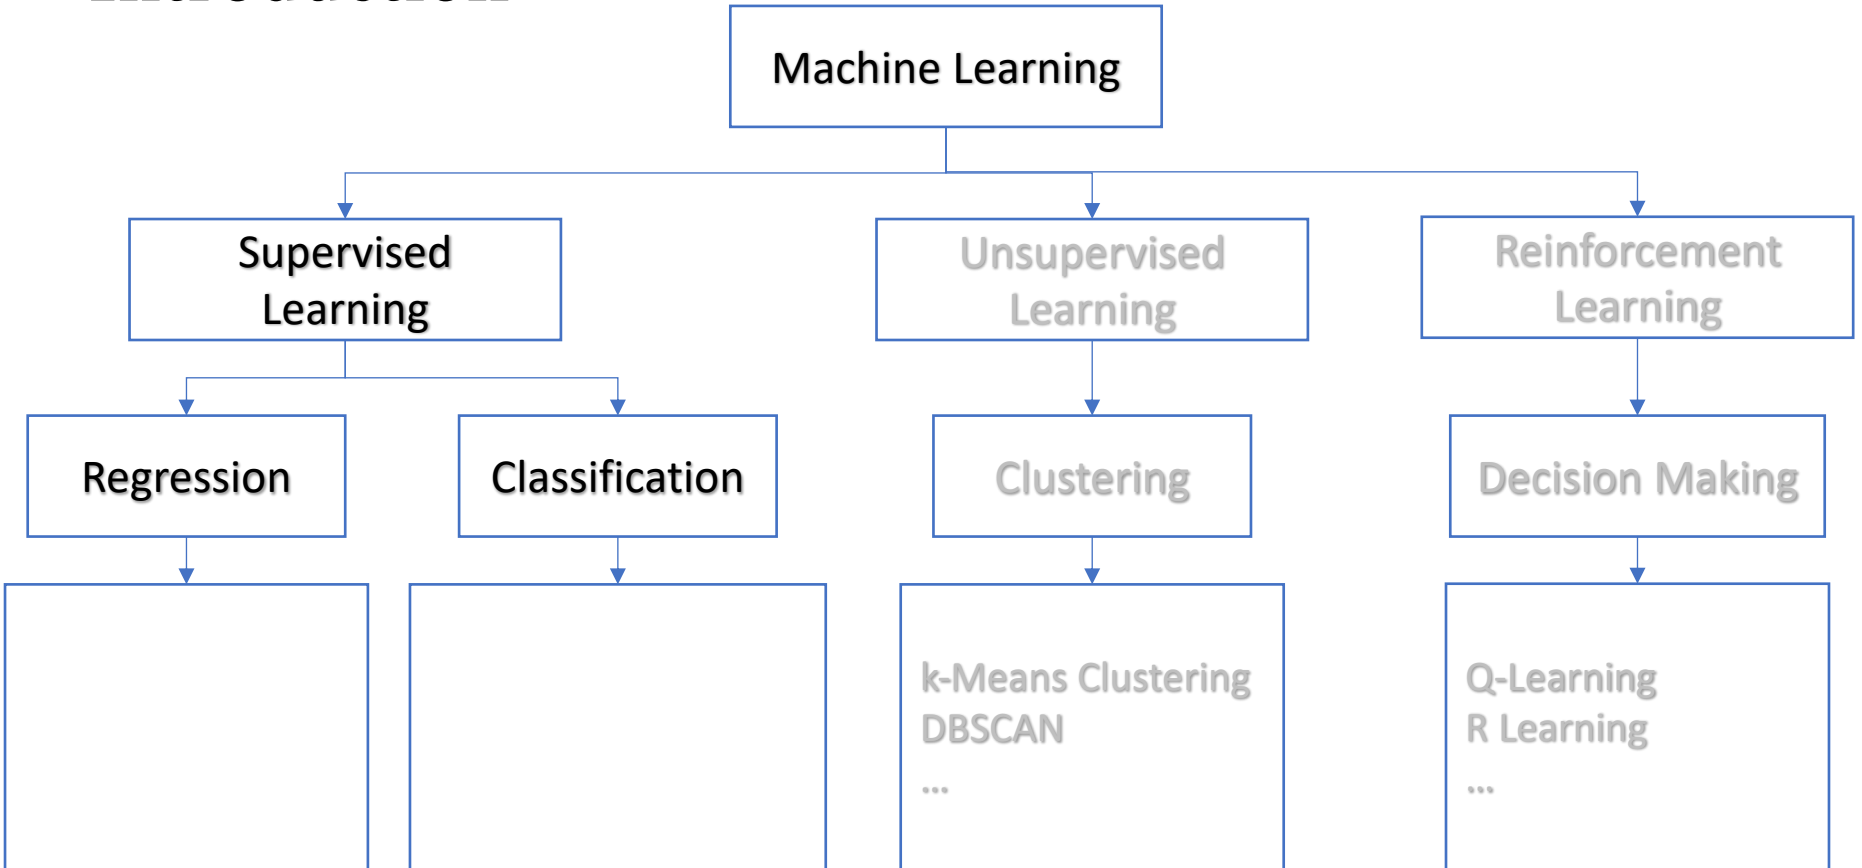
Overview

- **Introduction**
- k-Nearest Neighbours
- k-NN algorithm and pros/cons
- Evaluation Procedures

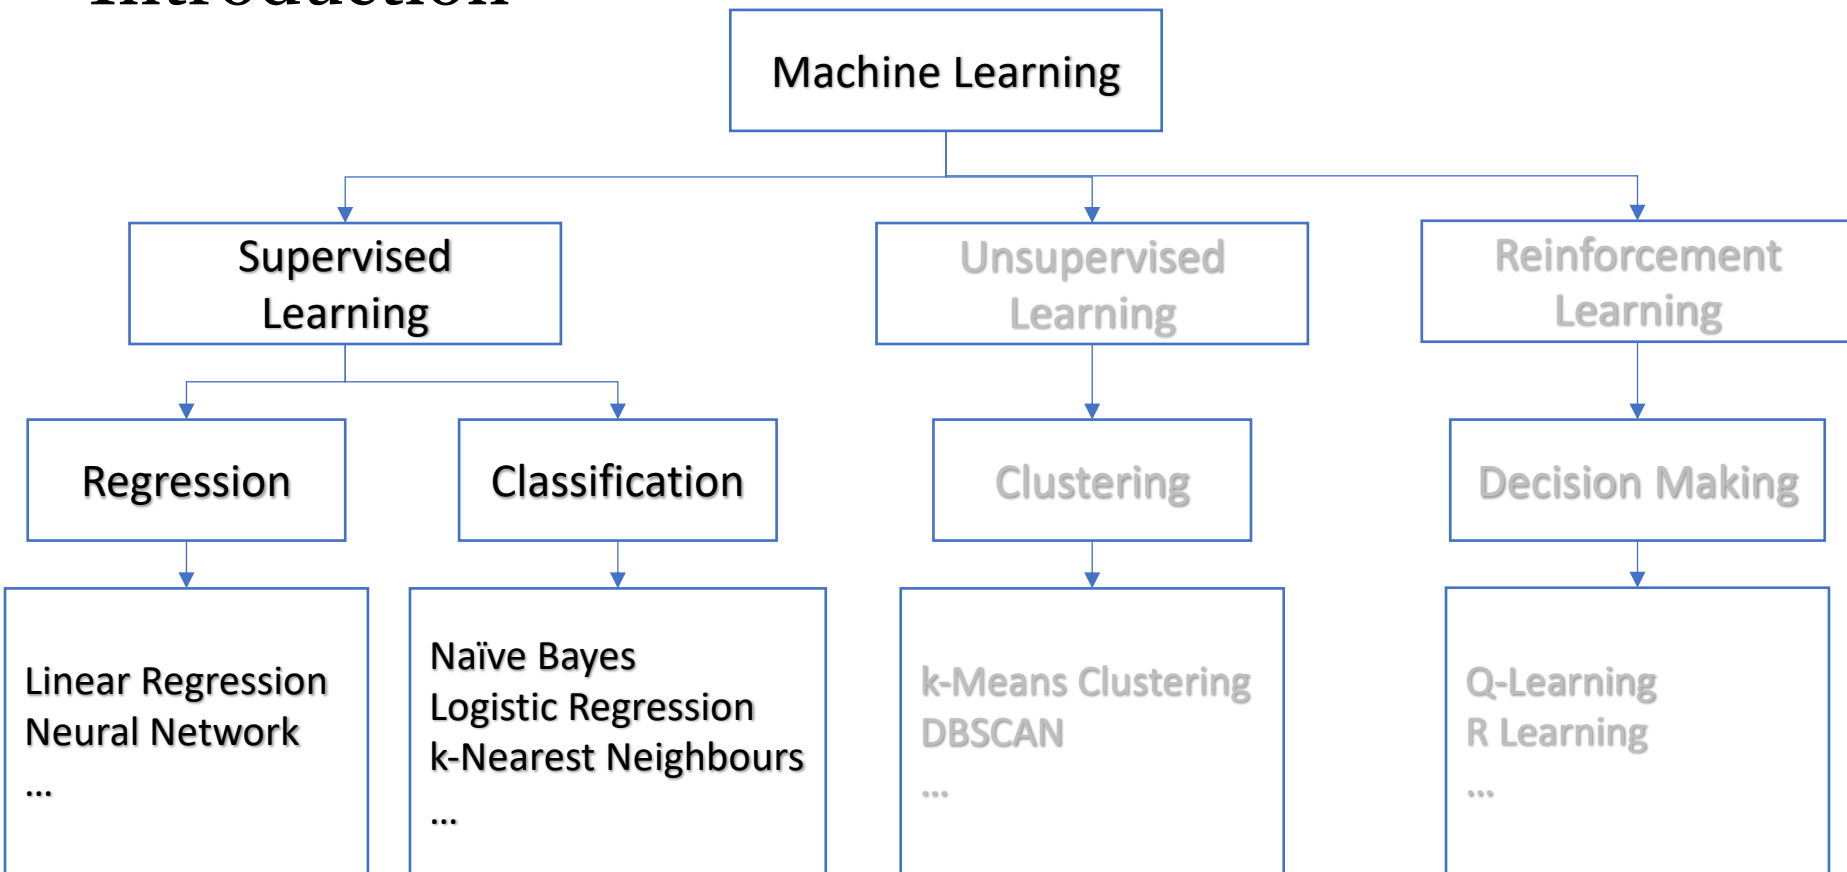
Introduction



Introduction



Introduction



Introduction

Machine Learning

Supervised
Learning

Regression

Linear Regression
Neural Network
...

Classification

Naïve Bayes
Logistic Regression
k-Nearest Neighbours
...

Unsupervised
Learning

Clustering

k-Means Clustering
DBSCAN
...

Reinforcement
Learning

Decision Making

Q-Learning
R Learning
...

Introduction

- Linear regression (as well as neural networks)
 - Uses the training data to estimate a fixed set of parameters \mathbf{w}
 - Defines our hypothesis $h_{\mathbf{w}}(\mathbf{x})$, thus making the training data useless
- A learning model that summarises data with a set of parameters of fixed size is called **parametric model**

Introduction

- Linear regression (as well as neural networks)
 - Uses the training data to estimate a fixed set of parameters \mathbf{w}
 - Defines our hypothesis $h_{\mathbf{w}}(\mathbf{x})$, thus making the training data useless
- A learning model that summarises data with a set of parameters of fixed size is called **parametric model**
- In this lecture, we will be using the following notation:
 - Variables are denoted by lowercase letters, e.g., x or y
 - Vectors are denoted by letters in bold, e.g., \mathbf{x}

Nonparametric Models

- A nonparametric model is a model that cannot be characterised by a bounded set of parameters
 - For instance, suppose that each prediction we make will consider all training examples, including the one from the previous prediction(s)
 - The set of examples grows over time, thus nonparametric

Nonparametric Models

- A nonparametric model is a model that cannot be characterised by a bounded set of parameters
 - For instance, suppose that each prediction we make will consider all training examples, including the one from the previous prediction(s)
 - The set of examples grows over time, thus nonparametric
- This approach is also called **instance-** or **memory-based learning**
 - The simplest method for instance-based learning is **table lookup**
 - For table lookup, we put all training examples in a table, and when looking for a value, we return the corresponding value
 - Problem: if the value does not exist, then a default value is returned

Nonparametric Models

- We can improve on table lookup with a slight variation:
 - Given a query \mathbf{x}_q , find the k examples that are *nearest* to \mathbf{x}_q
- This variation is called **k-Nearest Neighbours** (or **k-Nearest Neighbors**)
- Also, we can use the abbreviation k-NN
- The notation $NN(k, \mathbf{x}_q)$ denotes the set of k nearest neighbours

Overview

- Introduction
- **k-Nearest Neighbours**
- k-NN algorithm and pros/cons
- Evaluation Procedures

k-Nearest Neighbours (k-NN or KNN)

- Consider the following two-dimensional problem
 - Two classes: green or red ($y \in \{g, r\}$)
 - New example (blue) to classify (majority vote)




Image: taken from Ashraf *et al.*, "A Review of Intrusion Detection Systems Using Machine and Deep Learning in Internet of Things: Challenges, Solutions and Future Directions", *Electronics*, vol. 9, no. 7, 2020.

k-Nearest Neighbours (k-NN or KNN)

- Consider the following two-dimensional problem
 - Two classes: green or red ($y \in \{g, r\}$)
 - New example (blue) to classify (majority vote)
 - Look at the k nearest neighbours
 - Let $k = 3$, to avoid issues
 - We want to predict the class of the new example




Image: taken from Ashraf *et al.*, "A Review of Intrusion Detection Systems Using Machine and Deep Learning in Internet of Things: Challenges, Solutions and Future Directions", *Electronics*, vol. 9, no. 7, 2020.

k-Nearest Neighbours (k-NN or KNN)

- Consider the following two-dimensional problem

- Two classes: green or red ($y \in \{g, r\}$)
- New example (blue) to classify (majority vote)
- Look at the k nearest neighbours
- Let $k = 3$, to avoid issues
- We want to predict the class of the new example
- In this case, we predict green
- Intuitively, a distance metric on the **input space**




Image: taken from Ashraf *et al.*, "A Review of Intrusion Detection Systems Using Machine and Deep Learning in Internet of Things: Challenges, Solutions and Future Directions", *Electronics*, vol. 9, no. 7, 2020.

Distance Metrics

- Consider a problem with n dimensions, $\mathbf{x}^{(q)}$ being the new example
- The **Minkowski distance** (or L^p norm) is defined as

$$L^p(\mathbf{x}^{(q)}, \mathbf{x}^{(i)}) = \sqrt[p]{\sum_{j=1}^n |x_j^{(q)} - x_j^{(i)}|^p}$$

Distance Metrics

- With $p = 1$, the previous distance reduces to the **Manhattan distance**
- With Boolean attribute values, the so-called **Hamming distance** is used
- In general, the **Euclidean distance** is used, namely, when $p = 2$

$$L^2(\mathbf{x}^{(q)}, \mathbf{x}^{(i)}) = \sqrt{\sum_{j=1}^n (x_j^{(q)} - x_j^{(i)})^2}$$

Example

- Let us consider a problem with 2 dimensions (green and red), $\mathbf{x}^{(4)} = [0.1, 0.6]$ being the new example and the following training examples are given: $\mathbf{x}^{(1)} = [0.4, 0.2] \in \{green\}$, $\mathbf{x}^{(2)} = [0.4, 0.1] \in \{green\}$, $\mathbf{x}^{(3)} = [0.2, 0.6] \in \{red\}$
- Let us use the Euclidean distance to predict the class of $\mathbf{x}^{(4)}$, with $k = 1$

$$L^2(\mathbf{x}^{(4)}, \mathbf{x}^{(1)}) = \sqrt{\sum_{j=1}^n (x_j^{(4)} - x_j^{(1)})^2} =$$

Example

- Let us consider a problem with 2 dimensions (green and red), $\mathbf{x}^{(4)} = [0.1, 0.6]$ being the new example and the following training examples are given: $\mathbf{x}^{(1)} = [0.4, 0.2] \in \{green\}$, $\mathbf{x}^{(2)} = [0.4, 0.1] \in \{green\}$, $\mathbf{x}^{(3)} = [0.2, 0.6] \in \{red\}$
- Let us use the Euclidean distance to predict the class of $\mathbf{x}^{(4)}$, with $k = 1$

$$L^2(\mathbf{x}^{(4)}, \mathbf{x}^{(1)}) = \sqrt{\sum_{j=1}^n (x_j^{(4)} - x_j^{(1)})^2} = \sqrt{(x_1^{(4)} - x_1^{(1)})^2 + (x_2^{(4)} - x_2^{(1)})^2}$$

Example

- Let us consider a problem with 2 dimensions (green and red), $\mathbf{x}^{(4)} = [0.1, 0.6]$ being the new example and the following training examples are given: $\mathbf{x}^{(1)} = [0.4, 0.2] \in \{green\}$, $\mathbf{x}^{(2)} = [0.4, 0.1] \in \{green\}$, $\mathbf{x}^{(3)} = [0.2, 0.6] \in \{red\}$
- Let us use the Euclidean distance to predict the class of $\mathbf{x}^{(4)}$, with $k = 1$

$$\begin{aligned} L^2(\mathbf{x}^{(4)}, \mathbf{x}^{(1)}) &= \sqrt{\sum_{j=1}^n (x_j^{(4)} - x_j^{(1)})^2} = \sqrt{(x_1^{(4)} - x_1^{(1)})^2 + (x_2^{(4)} - x_2^{(1)})^2} \\ &= \sqrt{(0.1 - 0.4)^2 + (0.6 - 0.2)^2} = \sqrt{0.09 + 0.16} = 0.5 \end{aligned}$$

Example

- We repeat the process for all the points:

$$\begin{aligned} L^2(\mathbf{x}^{(4)}, \mathbf{x}^{(1)}) &= 0.5 \\ L^2(\mathbf{x}^{(4)}, \mathbf{x}^{(2)}) &= \\ L^2(\mathbf{x}^{(4)}, \mathbf{x}^{(3)}) &= \end{aligned}$$

Example

- We repeat the process for all the points:

$$\begin{aligned}L^2(\mathbf{x}^{(4)}, \mathbf{x}^{(1)}) &= 0.5 \\L^2(\mathbf{x}^{(4)}, \mathbf{x}^{(2)}) &= 0.583 \\L^2(\mathbf{x}^{(4)}, \mathbf{x}^{(3)}) &= 0.1\end{aligned}$$

Example

- We repeat the process for all the points:

$$\begin{aligned}L^2(\mathbf{x}^{(4)}, \mathbf{x}^{(1)}) &= 0.5 \\L^2(\mathbf{x}^{(4)}, \mathbf{x}^{(2)}) &= 0.583 \\L^2(\mathbf{x}^{(4)}, \mathbf{x}^{(3)}) &= 0.1\end{aligned}$$

Example

- We repeat the process for all the points:

$$\begin{aligned}L^2(\mathbf{x}^{(4)}, \mathbf{x}^{(1)}) &= 0.5 \\L^2(\mathbf{x}^{(4)}, \mathbf{x}^{(2)}) &= 0.583 \\L^2(\mathbf{x}^{(4)}, \mathbf{x}^{(3)}) &= 0.1\end{aligned}$$

- What happens if $k = 2$? What if $k = 3$?

Overview

- Introduction
- k-Nearest Neighbours
- **k-NN algorithm and pros/cons**
- Evaluation Procedures

k-NN Algorithm

Input: training examples $x^{(i)} \in \mathcal{X}$ and their corresponding class $y^{(i)}$, a new query example $x^{(q)}$, number of neighbours k

Output: prediction of the new query example $x^{(q)}$

For each training example $x^{(i)} \in \mathcal{X}$

 Calculate the distance between the training example $x^{(i)}$ and the new query example $x^{(q)}$

 Keep the best k distances (the shortest distance) in a data structure T

Return the majority (or plurality in case of non-binary classification) vote (or average/median) of the class $y^{(i)}$ for the first k entries of T

k-NN for Regression Problems

- Consider the following two-dimensional problem
 - Let us consider examples that take a value between 1 and 5 ($y \in [1, 5]$)
 - Each example has a value in y , e.g., 3.2 or 4.1

k-NN for Regression Problems

- Consider the following two-dimensional problem
 - Let us consider examples that take a value between 1 and 5 ($y \in [1, 5]$)
 - Each example has a value in y , e.g., 3.2 or 4.1
 - We calculate the distances between the new query example and all the other examples in the training set
 - Look at the k nearest neighbours
 - We predict the value of the new example
 - In this case, we predict the average or median of the values of the k nearest neighbours

Example

- Let us consider a problem with 2 dimensions (green and red), $\mathbf{x}^{(4)} = [0.1, 0.6]$ being the new example and the following training examples are given: $\mathbf{x}^{(1)} = [0.4, 0.2] \in \{3.6\}$, $\mathbf{x}^{(2)} = [0.4, 0.1] \in \{3.9\}$, $\mathbf{x}^{(3)} = [0.2, 0.6] \in \{2.2\}$
- Let $k = 2$. We know the distances from the previous example:
$$L^2(\mathbf{x}^{(4)}, \mathbf{x}^{(1)}) = 0.5$$
$$L^2(\mathbf{x}^{(4)}, \mathbf{x}^{(2)}) = 0.583$$
$$L^2(\mathbf{x}^{(4)}, \mathbf{x}^{(3)}) = 0.1$$

Example

- Let us consider a problem with 2 dimensions (green and red), $\mathbf{x}^{(4)} = [0.1, 0.6]$ being the new example and the following training examples are given: $\mathbf{x}^{(1)} = [0.4, 0.2] \in \{3.6\}$, $\mathbf{x}^{(2)} = [0.4, 0.1] \in \{3.9\}$, $\mathbf{x}^{(3)} = [0.2, 0.6] \in \{2.2\}$
- Let $k = 2$. We know the distances from the previous example:
- The average is: $y^{(4)} = 2.9$

Problem with Numeric Independent Variables

- Different numeric attributes may have different scales
- For example, if x_1 is in $[0,1]$ and x_2 is in $[1, 10]$, x_2 will affect the distance more

Problem with Numeric Independent Variables

- To avoid this problem, we normalise the numeric input attributes of all data as in the following

$$\text{normalise} \left(x_j^{(i)} \right) = \frac{x_j^{(i)} - \min_j}{\max_j - \min_j}$$

- Another approach is to calculate mean μ_j and standard deviation σ_j for each dimension j as: $(x_j^{(i)} - \mu_j) / \sigma_j$

Example

- Consider examples with two dimensions, where x_1 represents the age of a patient and x_2 represents their weight, $y \in \{yes, no\}$
- Let us calculate the normalised values for $x^{(1)}$:

Days	x_1	x_2	y
$x^{(1)}$	14	70	yes
$x^{(2)}$	12	90	no
$x^{(3)}$	15	66	yes

Example

- Consider examples with two dimensions, where x_1 represents the age of a patient and x_2 represents their weight, $y \in \{yes, no\}$
- Let us calculate the normalised values for $\mathbf{x}^{(1)}$:
- $normalise(x_1^{(1)}) = \frac{x_1^{(1)} - \min_1}{\max_1 - \min_1} =$
- $normalise(x_2^{(1)}) = \frac{x_2^{(1)} - \min_2}{\max_2 - \min_2} =$

Days	x_1	x_2	y
$\mathbf{x}^{(1)}$	14	70	yes
$\mathbf{x}^{(2)}$	12	90	no
$\mathbf{x}^{(3)}$	15	66	yes

Example

- Consider examples with two dimensions, where x_1 represents the age of a patient and x_2 represents their weight, $y \in \{yes, no\}$
- Let us calculate the normalised values for $\mathbf{x}^{(1)}$:
- $normalise(x_1^{(1)}) = \frac{x_1^{(1)} - \min_1}{\max_1 - \min_1} = \frac{14 - 12}{15 - 12} = 0.667$
- $normalise(x_2^{(1)}) = \frac{x_2^{(1)} - \min_2}{\max_2 - \min_2} = \frac{70 - 66}{90 - 66} = 0.167$

Days	x_1	x_2	y
$\mathbf{x}^{(1)}$	14	70	yes
$\mathbf{x}^{(2)}$	12	90	no
$\mathbf{x}^{(3)}$	15	66	yes

k-NN Algorithm with Normalisation

Input: training examples $x^{(i)} \in \mathcal{X}$ and their corresponding class $y^{(i)}$, a new query example $x^{(q)}$, number of neighbours k

Output: prediction of the new query example $x^{(q)}$

For each training example $x^{(i)} \in \mathcal{X}$

Calculate the **normalised** distance between the training example $x^{(i)}$ and the new query example $x^{(q)}$

Keep the best k distances (the shortest distance) in a data structure T

Return the majority (or plurality in case of non-binary classification) vote (or average/median) of the class $y^{(i)}$ for the first k entries of T

Different Input Attributes

- For **numeric** input attributes, e.g., age in $[0, 100]$, we calculate the distance as shown in previous examples
- For **ordinal** input attributes, e.g., sunny in $\{\text{yes}, \text{no}\}$, we can convert the values to numeric values: yes = 1, no = 0
- For **categorical** input attributes, e.g., phone_brand in $\{\text{samsung}, \text{apple}, \text{nokia}\}$, we can use the following approach:
 - If the value of the query example is the same as the value for example i , then their difference is 0. Formally, if $(x_j^{(q)} = x_j^{(i)})$, then $(x_j^{(q)} - x_j^{(i)}) = 0$
 - Otherwise, their difference is 1. Formally, if $(x_j^{(q)} \neq x_j^{(i)})$, then $(x_j^{(q)} - x_j^{(i)}) = 1$

Summary

k-NN Learning Algorithm

- The algorithm does not have proper training
- We simply store all training data, which increase over time
- We normalise by calculating the minimum and maximum in the training data

k-NN Model

- All training data, the values of the numeric input attributes

k-NN prediction for an instance ($\mathbf{x}^{(q)}$, $y = ?$)

- Find the k nearest neighbours whose distance to $\mathbf{x}^{(i)}$ is the smallest
- For classification problems, majority vote. For regression problems, average/median

Pros and Cons of k-NN

Pros

- Training is simple and fast: just store training data
- Find the class of the new example based on most similar examples present in the training data

Cons

- It uses large space in memory: we need to store all data
- Running the algorithm can be slow if we have many training examples and many dimensions

Complexity

- Given a set of N examples and a query example, we calculate the distance to the query example from each one and keep the best k


Complexity

- Given a set of N examples and a query example, we calculate the distance to the query example from each one and keep the best k
- The above runs in $O(N)$ with a (sequential) table
- However, these methods are designed for large datasets. Therefore, better data structure are desirable


Complexity

- Given a set of N examples and a query example, we calculate the distance to the query example from each one and keep the best k
- The above runs in $O(N)$ with a (sequential) table
- However, these methods are designed for large datasets. Therefore, better data structure are desirable
- A binary tree runs in $O(\log N)$, whereas a hash table runs in $O(1)$

Discussion




($k=1$)




($k=5$)

- Nonparametric methods are still subject to under- and overfitting

Discussion



($k=1$)



($k=5$)

- Nonparametric methods are still subject to under- and overfitting
- In the above case, 1-NN is overfitting as it reacts too much to the black outlier in the upper right and the white at (5.4, 3.7)
- The 5-NN decision boundary is good; higher k would underfit

Overview

- Introduction
- k-Nearest Neighbours
- k-NN algorithm and pros/cons
- **Evaluation Procedures**


Evaluation Procedures

- A supervised learning method consists of 3 main elements:
 - Model, namely, the form of function we want to learn (with free parameters)
 - Cost function, namely, the misfit between a particular function from the model (provided a training set)
 - Training algorithm, namely, gradient descent minimisation of cost function
- Running the training algorithm on training data learns the “best” values of the free parameters, yielding a predictor


This part is adapted from Ata Kaban's slides and worked examples from Andrew Moore's tutorial slides
<https://sites.astro.caltech.edu/~george/aybi199/AMooreTutorials/>

Evaluation of Predictor before Deployment

- We use evaluation procedures to determine how good our model is. In other words, we want to estimate the future performance of a predictor
- To do so, we randomly split the available annotated data into:
 - A **training set** – to estimate all the free parameters
 - A **test set** – to evaluate the trained predictor before deployment



Choosing the Model



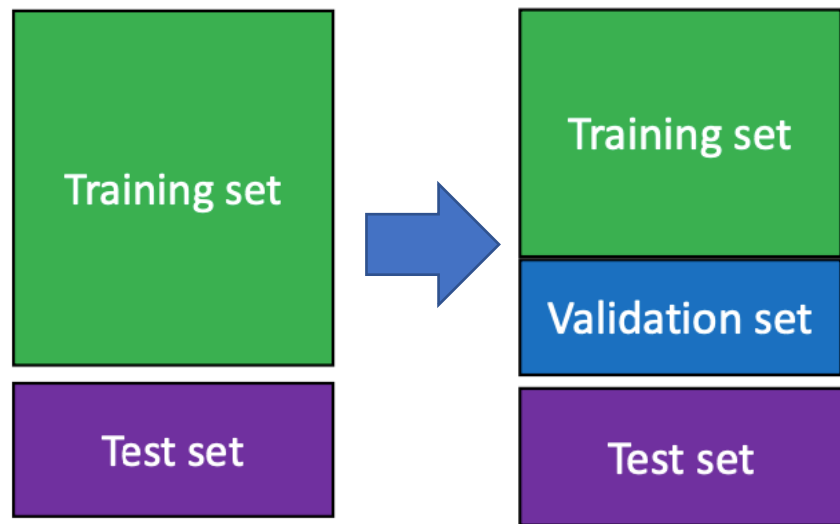
- Hyperparameters are “higher-level” free parameters
- Each hyperparameter value corresponds to a different model
- Which model? We need a **criterion to estimate future performance**

Evaluating Models for Model Choice

- Do not confuse this with evaluating a predictor (model already chosen)
- The training set is used for training within a chosen model
- The test set is used to evaluate the performance of the trained predictor
- **None of the above can be used to choose the model!**
 - If we are tempted to use the test set, we no longer have an independent dataset to evaluate the final predictor before deployment


Evaluating Models for Model Choice

- Idea: to choose between models or hyperparameters, take a subset of the training set and create a validation set
- Methods
 - Holdout validation
 - Cross-validation
 - Leave-one-out validation




Method 1: Holdout Validation


1. Randomly choose 30% of data to form a **validation set**
2. Keep the rest in the **training set**
3. **Train your model on the training set**
4. Estimate the test performance on the **validation set**
5. **Choose model with lowest validation error**
6. Re-train with chosen model on **joined training and validation** to obtain predictor
7. Estimate future performance on test set
8. Ready to deploy predictor




Example: Holdout Validation



Model 1
Mean Squared Validation Error = 2.4




Model 2
Mean Squared Validation Error = 0.9




Model 3
Mean Squared Validation Error = 2.2

Example: Holdout Validation




Model 1

Mean Squared Validation Error = 2.4



Model 2

Mean Squared Validation Error = 0.9



Model 3

Mean Squared Validation Error = 2.2

Method 1: Holdout Validation, Step 4

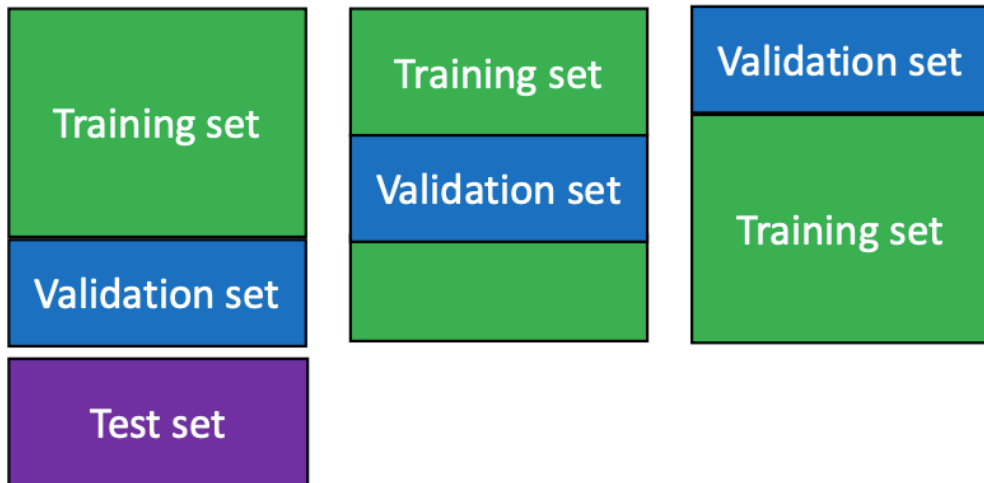
4. Estimate the test performance on the validation set

- Different approaches for regression and for classification:
 - In regression, we compute the cost function (mean square error) on the examples of the validation set (instead of the training set)
 - In classification, we compute the 0-1 error metric on validation set, i.e.,

$$\frac{\text{\# wrong predictions}}{\text{\# predictions}} = 1 - \text{Accuracy}$$


Method 2: k-Fold Cross-Validation

1. Randomly split the training set into k disjoint sets of equal size ($k=3$ in this example)
2. Use the $k-1$ of those together for training
3. Use the remaining one for validation
4. Permute the k sets and repeat k times
5. Average the performance on the k validation sets



Example: k-Fold Cross-Validation


Randomly split the dataset into $k=3$ partitions



Example: k-Fold Cross-Validation

Randomly split the dataset into $k=3$ partitions

Blue partition: train on all the data except the blue partition. Compute the validation error using the data in the blue partition




Example: k-Fold Cross-Validation

Randomly split the dataset into $k=3$ partitions

Blue partition: train on all the data except the blue partition. Compute the validation error using the data in the blue partition

Green partition: train on all the data except the green partition. Compute the validation error using the data in the green partition




Example: k-Fold Cross-Validation

Randomly split the dataset into $k=3$ partitions

Blue partition: train on all the data except the blue partition. Compute the validation error using the data in the blue partition


Green partition: train on all the data except the green partition. Compute the validation error using the data in the green partition

Purple partition: train on all the data except the purple partition. Compute the validation error using the data in the purple partition




Example: k-Fold Cross-Validation


Take the mean of these errors



Model 1
 $MSE_{3FOLD}=2.05$



Model 2
 $MSE_{3FOLD}=1.11$



Model 3
 $MSE_{3FOLD}=2.93$

Method 3: Leave-One-Out Validation


- We leave out a **single example** for validation, and train on all the remaining annotated data
- For a total of N examples, we repeat this process N times, each time leaving out a single example
- Take the average of the validation errors as measured on the left-out examples

Method 3: Leave-One-Out Validation

- We leave out a **single example** for validation, and train on all the remaining annotated data
- For a total of N examples, we repeat this process N times, each time leaving out a single example
- Take the average of the validation errors as measured on the left-out examples
- Same as N -fold cross-validation where N is the number of labelled examples

Advantages and Disadvantages

	Advantages	Disadvantages	
Holdout validation	Computationally cheapest	Most unreliable if sample size is not large enough	Large sample
3-fold	Slightly more reliable than holdout	<ul style="list-style-type: none">• Wastes 1/3-rd annotated data.• Computationally 3-times as expensive as holdout	
10-fold	<ul style="list-style-type: none">• Only wastes 10%• Fairly reliable	<ul style="list-style-type: none">• Wastes 10% annotated data• Computationally 10-times as expensive as holdout	
Leave-one-out	Doesn't waste data	Computationally most expensive	Small sample



Aims of the Session

You should now be able to:

- Explain the steps of k-Nearest Neighbours
- Apply k-Nearest Neighbours to problems involving numeric, ordinal and categorical input attributes
- Use evaluation procedures to estimate the performance of a predictor

References

- Russell, A. S., and Norvig, P. (2010), *Artificial Intelligence A Modern Approach*, 3rd Edition. Prentice Hall.
 - Chapter 18 – Learning from Examples (Section 18.8 up to 18.8.1, Section 18.8.4)