


# Dynamic Memory Management

Eike Ritter and Aad van Moorsel  
School of Computer Science  
University of Birmingham

# Recap: Memory layout of C program

- Local variables use Stack 'temporarily'. They are active only within their functions.
- Static variables are stored in the Data segments. They preserve their values.
- Global variables are stored in the Data segments. They are accessible to all functions of the C program.

[ This figure is symbolic. Heap section need not start from the address just after data segment ]



Use of Heap will be covered in Dynamic Memory Management.

# Application scenario


Consider an application where the volume of data is not known beforehand

Example: Store info of the people you meet during office hour

- The number of people is not known
- Each person may provide different amount of info

Solution 1: Allocate a large array

```
char info[10000][2000];
```



# Application scenario

Consider an application where the volume of data is not known beforehand

Example: Store info of the people you meet during office hour


- The number of people is not known
- Each person may provide different amount of info

Solution 1: Allocate a large array

```
char info[10000][2000];
```

Problems:

1. May be wasteful
2. May be insufficient
3. Tiny computes can't afford large amount of space



## Application scenario

Consider an application where the volume of data is not known beforehand

Example: Store info of the people you meet during office hour

- The number of people is not known
- Each person may provide different amount of info

**Best solution:** Allocate memory at **runtime** as per **demand**

Advantages: Optimal memory consumption depending on the current state of the application

# Dynamic memory allocation in C

# Dynamic memory allocation functions

C provides dynamic memory management functions in `stdlib.h`

- `malloc()`
  1. Allocates requested number of bytes of **contiguous** memory from the **Heap**.
  2. Returns a pointer to the first byte of the allocated space.
  3. If memory allocation fails, then `malloc()` returns NULL pointer.

Syntax:

```
T *p;  
p = (T *) malloc(number of bytes);
```

# Dynamic memory allocation functions

C provides dynamic memory management functions in `stdlib.h`

- `malloc()`
  1. Allocates requested number of bytes of **contiguous** memory from the **Heap**.
  2. Returns a pointer to the first byte of the allocated space.
  3. If memory allocation fails, then `malloc()` returns NULL pointer.

Example: Allocate space for 3 `int`

```
int *p;  
p = (int *) malloc(3*4); // int is 4  
bytes
```



Allocates 12 bytes of contiguous memory in Heap.  
p points to the first byte



# Dynamic memory allocation functions

C provides dynamic memory management functions in `stdlib.h`

- `malloc()`
  1. Allocates requested number of bytes of **contiguous** memory from the **Heap**.
  2. Returns a pointer to the first byte of the allocated space.
  3. If memory allocation fails, then `malloc()` returns NULL pointer.

Example: Allocate space for 3 `int`

```
int *p;  
p = (int *) malloc(3*sizeof(int));  
// Allocates memory for 3 integers
```

Use `sizeof()` for convenience.

It returns the size of a `data_type` in bytes.

# Dynamic memory allocation functions

C provides dynamic memory management functions in `stdlib.h`

- `malloc()`
  1. Allocates requested number of bytes of **contiguous** memory from the **Heap**.
  2. Returns a pointer to the first byte of the allocated space.
  3. If memory allocation fails, then `malloc()` returns NULL pointer.

Example: Allocate space for 3 `int`

```
int *p;  
if((p= (int *) malloc(3*sizeof(int)))==NULL){  
    printf("Allocation failed");  
    exit(-1);  
}
```

Good practice: check if a NULL pointer is returned by `malloc()`

## Releasing allocated memory using free()

- Dynamically allocated block of memory can be released using `free()`.
- After `free()`, the block of memory is returned to the Heap.

```
int *p;  
// Block of memory is allocated  
if((p= (int *) malloc(3*sizeof(int)))==NULL){  
    printf("Allocation failed");  
    exit(-1);  
}  
  
// [Some statements involving p]  
  
// Block of memory pointed by p is released  
free(p);
```

## Example: Array of variable length

```
#include<stdio.h>
#include<stdlib.h>
int main(){
    int N, i;
    int *p;
    printf("Provide array size:");
    scanf("%d", &N);
    // Allocate memory in Heap for
    // array pointed by p
    if((p= (int *) malloc(N*sizeof(int)))==NULL){
        printf("Allocation failed");
        exit(-1);
    }
    printf("Provide %d integers\n", N);
    for(i=0; i<N; i++)
        scanf("%d", p+i);

    free(p);
    return 0;
}
```

# Memory leak

If memory allocated using `malloc()` is not `free()`-ed, then the system will “leak memory”

- Block of memory is allocated, but not returned to Heap
- The program therefore grows larger over time
- In C, it is **your responsibility** to prevent memory leak

## Example: Memory leak

```
int main(){
    int N, i, *p;

    while(1){
        printf("Provide array size:");
        scanf("%d", &N);
        if((p= (int *) malloc(N*sizeof(int)))==NULL){
            printf("Allocation failed");
            exit(-1);
        }
        printf("Provide %d integers\n", N);
        for(i=0; i<N; i++)
            scanf("%d", p+i);

        //free(p);
    }
    return 0;
}
```

Each iteration of while loop allocates memory. But that memory is never freed. This causes memory leak.

# Valgrind: a tool for memory leak detection

- We will use the valgrind tool to detect memory leaks
- Quick start info: <http://valgrind.org/docs/manual/quick-start.html>

Steps to follow:

1. Compile your C code using gcc
2. Then use valgrind to run the compiled code:  
`valgrind --leak-check=full ./a.out`

## Example: valgrind output

```
int main()
{
    int *pt = malloc(1*sizeof(int));
    scanf("%d", pt);
    printf("%d\n", *pt);
    return 0;
}
```

Output of `valgrind --leak-check=full ./a.out`

```
==1057== Memcheck, a memory error detector
==1057== Copyright (C) 2002-2017, and GNU GPL'd, by Julian Seward et al.
==1057== Using Valgrind-3.13.0 and LibVEX; rerun with -h for copyright info
==1057== Command: ./a.out
==1057==
```

... continued

```
==1057== HEAP SUMMARY:
==1057==    in use at exit: 4 bytes in 1 blocks
==1057==    total heap usage: 1 allocs, 0 frees, 4 bytes allocated
==1057==
==1057== LEAK SUMMARY:
==1057==    definitely lost: 0 bytes in 0 blocks
==1057==    indirectly lost: 0 bytes in 0 blocks
==1057==    possibly lost: 0 bytes in 0 blocks
==1057==    still reachable: 4 bytes in 1 blocks
==1057==    suppressed: 0 bytes in 0 blocks
```

Valgrind reports a memory leak of 4 bytes. ■■



## Example: valgrind output

```
int main()
{
    int *pt = malloc(1*sizeof(int));
    scanf("%d", pt);
    printf("%d\n", *pt);
    free(pt);
    return 0;
}
```

→ Allocated memory is freed here

Output of `valgrind --leak-check=full ./a.out`

```
==2308== Memcheck, a memory error detector
==2308== Copyright (C) 2002-2017, and GNU GPL'd, by Julian Seward et al.
==2308== Using Valgrind-3.13.0 and LibVEX; rerun with -h for copyright info
==2308== Command: ./a.out
==2308==
5
5
==2308==
==2308== HEAP SUMMARY:
==2308==    in use at exit: 0 bytes in 0 blocks
==2308==    total heap usage: 1 allocs, 1 frees, 4 bytes allocated
==2308==
==2308== All heap blocks were freed -- no leaks are possible
==2308==
```

No memory leak reported! ◀◀


## Consequences of memory leak

- Memory leaks slowdown system performance by reducing the amount of available memory.
- In modern operating systems, memory used by an application is released when the application terminates
  - If a program (with leak) runs for a short duration, then no serious problem
  - Will be a serious problem when a leaking program runs for long duration
- Memory leak will cause serious issues on resource-constrained embedded devices

**Overall, your program should not contain memory leaks**

## Another memory leak example(1)

```
int main(){
    int *p1, **p2;
    p1 = malloc(sizeof(int));
    *p1=7;
    p2 = malloc(sizeof(int*));
    *p2=p1;
    return 0;
}
```




p2 is a pointer to a pointer variable

```
==3600== HEAP SUMMARY:
==3600==      in use at exit: 12 bytes in 2 blocks
==3600==    total heap usage: 2 allocs, 0 frees, 12 bytes allocated
==3600==
==3600== 12 (8 direct, 4 indirect) bytes in 1 blocks are definitely lost in loss record 2 of 2
==3600==    at 0x4C29BC3: malloc (vg_replace_malloc.c:299)
==3600==    by 0x400546: main (in /users/cosic/ssinharo/dynamicmem/a.out)
==3600==
==3600== LEAK SUMMARY:
==3600==    definitely lost: 8 bytes in 1 blocks
==3600==    indirectly lost: 4 bytes in 1 blocks
==3600==    possibly lost: 0 bytes in 0 blocks
==3600==    still reachable: 0 bytes in 0 blocks
==3600==         suppressed: 0 bytes in 0 blocks
==3600==
```

## Another memory leak example(2)

```
int main(){
    int *p1, **p2;
    p1 = malloc(sizeof(int));
    *p1=7;
    p2 = malloc(sizeof(int*));
    *p2=p1;
    free(p1);
    return 0;
}
```




p2 is a pointer to a pointer variable

```
==3847== LEAK SUMMARY:
==3847==    definitely lost: 8 bytes in 1 blocks
==3847==    indirectly lost: 0 bytes in 0 blocks
==3847==    possibly lost: 0 bytes in 0 blocks
==3847==    still reachable: 0 bytes in 0 blocks
==3847==    suppressed: 0 bytes in 0 blocks
```

Still 8 bytes are leaked!

## Another memory leak example(3)

```
int main(){
    int *p1, **p2;
    p1 = malloc(sizeof(int));
    *p1=7;
    p2 = malloc(sizeof(int*));
    *p2=p1;
    free(p2);
    return 0;
}
```




p2 is a pointer to a pointer variable

```
==6994== LEAK SUMMARY:
==6994==      definitely lost: 4 bytes in 1 blocks
==6994==      indirectly lost: 0 bytes in 0 blocks
==6994==      possibly lost: 0 bytes in 0 blocks
==6994==      still reachable: 0 bytes in 0 blocks
==6994==      suppressed: 0 bytes in 0 blocks
```

4 bytes are leaked!

## Another memory leak example(4)

```
int main(){
    int *p1, **p2;
    p1 = malloc(sizeof(int));
    *p1=7;
    p2 = malloc(sizeof(int*));
    *p2=p1;
    free(p1);
    free(p2);
    return 0;
}
```



p2 is a pointer to a pointer variable

No memory leak!

## Accessing memory after free( )

- After `free(p)`, the memory is no longer owned by the program.



- Some other program may be using the memory! **Data is lost**

- Re-accessing** the memory pointed by `p` causes **undefined behaviour**

```
int main(){
    int sum, *p;
    p = (int *) malloc(4*sizeof(int));
    // [some code here]
    free(p);
    sum = sum + *p; //Use after free issue
    // [some code here]
}
```

Memory pointed by `p` was freed and then used.  
Program will cause undefined behaviour.  
Valgrind reports memory error.

# Don't free something that did not come from malloc()

```
foo(int a){  
    int b;  
    int c[]={1, 2, 3};  
    free(&a);  
    free(&b);  
    free(c);  
}
```

- Only, dynamically created objects are stored in heap
- gcc gives warnings: “attempt to free a non-heap object”
- However, program crashes with segmentation fault

Valgrind reports memory errors



## Double free problem

- Double free errors occur when `free()` is called more than once with the same memory address as an argument.

```
int main(){  
    char *p1=malloc(1);  
    char *p2=malloc(1);  
    free(p1);  
    free(p1);  
}
```

- When `free()` is called twice with the same argument, the program's memory management data structures become corrupted
- Consequences: Program malfunction including security vulnerabilities

At runtime, modern OS detects double-free operation.

**Program is aborted.**

# How to think about double free

```
int main(){  
    char *p1=malloc(1);  
    char *p2=malloc(1);  
    free(p1);  
    free(p1);  
}
```

Heap manager

## Available blocks

...
...
p1
p1

- Every time `free()` is called, the address is added in the list of available memory blocks
- The last address added to the list can be by the next `malloc()` call
- Since p1 is present twice in the list, next two memory allocations will have the same address (**which is a problem**)
- Valgrind reports memory error